

INCOME QUALIFICATION

February 7, 2023

INCOME QUALIFICATION PROJECT-2

SUBMITTED BY-SOMYA KUMARI PANDEY

Problem Statement Scenario: Many social programs have a hard time making sure the right people are given enough aid. It's tricky when a program focuses on the poorest segment of the population. This segment of population can't provide the necessary income and expense records to prove that they qualify.

In Latin America, a popular method called Proxy Means Test (PMT) uses an algorithm to verify income qualification. With PMT, agencies use a model that considers a family's observable household attributes like the material of their walls and ceiling or the assets found in their homes to classify them and predict their level of need. While this is an improvement, accuracy remains a problem as the region's population grows and poverty declines.

The Inter-American Development Bank (IDB) believes that new methods beyond traditional econometrics, based on a dataset of Costa Rican household characteristics, might help improve PMT's performance.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[2]: train=pd.read_csv('train.csv')
test=pd.read_csv('test.csv')
```

Let us explore our dataset before moving further

```
[3]: print('Shape of train dataset is {}'.format(train.shape))
print('Shape of test dataset is {}'.format(test.shape))
```

Shape of train dataset is (9557, 143)

Shape of test dataset is (23856, 142)

Let us identify our target variable

```
[4]: for i in train.columns:
    if i not in test.columns:
        print("Our Target variable is {}".format(i))
```

Our Target variable is Target

Lets Understand the type of data.

```
[5]: print(train.dtypes.value_counts())
```

```
int64      130
float64      8
object       5
dtype: int64
```

```
[6]: print(train.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9557 entries, 0 to 9556
Columns: 143 entries, Id to Target
dtypes: float64(8), int64(130), object(5)
memory usage: 10.4+ MB
None
```

We have mixed data types. Specified as below:**

- float64 : 8 variables
- int64 : 130 variables
- object :5 variables

```
[7]: #lets explore each different types of datasets
for i in train.columns:
    a=train[i].dtype
    if a == 'object':
        print(i)
```

```
Id
idhogar
dependency
edjefe
edjefa
```

Below is Data dictionary for above object variables

- ID = Unique ID
- idhogar, Household level identifier
- dependency, Dependency rate, calculated = (number of members of the household younger than 19 or older than 64)/(number of member of household between 19 and 64)
- edjefe, years of education of male head of household, based on the interaction of escolar (years of education), head of household and gender, yes=1 and no=0
- edjefa, years of education of female head of household, based on the interaction of escolar (years of education), head of household and gender, yes=1 and no=0

```
[8]: # lets drop Id variable.

train.drop(['Id', 'idhogar'], axis=1, inplace=True)
```

```
[9]: train['dependency'].value_counts()
```

```
[9]: yes          2192
no            1747
.5           1497
2             730
1.5           713
.33333334     598
.66666669     487
8             378
.25           260
3             236
4             100
.75            98
.2            90
1.3333334     84
.40000001     84
2.5            77
5             24
.80000001     18
1.25          18
3.5           18
2.25          13
.71428573     12
1.75          11
.83333331     11
1.2           11
.22222222     11
.2857143       9
1.6666666     8
.60000002     8
.16666667     7
6              7
Name: dependency, dtype: int64
```

Lets Convert object variables into numerical data

```
[10]: def map(i):

        if i=='yes':
            return(float(1))
        elif i=='no':
            return(float(0))
```

```

else:
    return(float(i))

```

```
[11]: train['dependency']=train['dependency'].apply(map)
```

```
[12]: for i in train.columns:
        a=train[i].dtype
        if a == 'object':
            print(i)

```

edjefe
edjefa

```
[13]: train.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9557 entries, 0 to 9556
Columns: 141 entries, v2a1 to Target
dtypes: float64(9), int64(130), object(2)
memory usage: 10.3+ MB

```

```
[14]: train['edjefe']=train['edjefe'].apply(map)
        train['edjefa']=train['edjefa'].apply(map)

```

```
[15]: train.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9557 entries, 0 to 9556
Columns: 141 entries, v2a1 to Target
dtypes: float64(11), int64(130)
memory usage: 10.3 MB

```

Now all data is in numerical form Lets identify variable with 0 variance

```
[16]: var_df=pd.DataFrame(np.var(train,0),columns=['variance'])
        var_df.sort_values(by='variance').head(15)
        print('Below are columns with variance 0.')
        col=list((var_df[var_df['variance']==0]).index)
        print(col)

```

Below are columns with variance 0.
['elimbasu5']

elimbasu5 : 1 if rubbish disposal mainly by throwing in river, creek or sea.

Interpretation :From above it is shown that all values of elimbasu5 is same so there is no variability in dataset therefor we will drop this variable

Check if there are any biases in your dataset

```
[17]: contingency_tab=pd.crosstab(train['r4t3'],train['hogar_total'])
Observed_Values=contingency_tab.values
import scipy.stats
b=scipy.stats.chi2_contingency(contingency_tab)
Expected_Values = b[3]
no_of_rows=len(contingency_tab.iloc[0:2,0])
no_of_columns=len(contingency_tab.iloc[0,0:2])
df=(no_of_rows-1)*(no_of_columns-1)
print("Degree of Freedom:-",df)
from scipy.stats import chi2
chi_square=sum([(o-e)**2./e for o,e in zip(Observed_Values,Expected_Values)])
chi_square_statistic=chi_square[0]+chi_square[1]
print("chi-square statistic:-",chi_square_statistic)
alpha=0.05
critical_value=chi2.ppf(q=1-alpha,df=df)
print('critical_value:',critical_value)
p_value=1-chi2.cdf(x=chi_square_statistic,df=df)
print('p-value:',p_value)
print('Significance level: ',alpha)
print('Degree of Freedom: ',df)
print('chi-square statistic:',chi_square_statistic)
print('critical_value:',critical_value)
print('p-value:',p_value)
if chi_square_statistic>=critical_value:
    print("Reject H0,There is a relationship between 2 categorical variables")
else:
    print("Retain H0,There is no relationship between 2 categorical variables")

if p_value<=alpha:
    print("Reject H0,There is a relationship between 2 categorical variables")
else:
    print("Retain H0,There is no relationship between 2 categorical variables")
```

```
Degree of Freedom:- 1
chi-square statistic:- 17022.072400560897
critical_value: 3.841458820694124
p-value: 0.0
Significance level: 0.05
Degree of Freedom: 1
chi-square statistic: 17022.072400560897
critical_value: 3.841458820694124
p-value: 0.0
Reject H0,There is a relationship between 2 categorical variables
Reject H0,There is a relationship between 2 categorical variables
```

Therefore,variables ('r4t3','hogar_total') have relationship between them. For good result we can use any one of them.

```
[18]: contingency_tab=pd.crosstab(train['tipovivi3'],train['v2a1'])
Observed_Values=contingency_tab.values
import scipy.stats
b=scipy.stats.chi2_contingency(contingency_tab)
Expected_Values = b[3]
no_of_rows=len(contingency_tab.iloc[0:2,0])
no_of_columns=len(contingency_tab.iloc[0,0:2])
df=(no_of_rows-1)*(no_of_columns-1)
print("Degree of Freedom:-",df)
from scipy.stats import chi2
chi_square=sum([(o-e)**2./e for o,e in zip(Observed_Values,Expected_Values)])
chi_square_statistic=chi_square[0]+chi_square[1]
print("chi-square statistic:-",chi_square_statistic)
alpha=0.05
critical_value=chi2.ppf(q=1-alpha,df=df)
print('critical_value:',critical_value)
p_value=1-chi2.cdf(x=chi_square_statistic,df=df)
print('p-value:',p_value)
print('Significance level: ',alpha)
print('Degree of Freedom: ',df)
print('chi-square statistic:',chi_square_statistic)
print('critical_value:',critical_value)
print('p-value:',p_value)
if chi_square_statistic>=critical_value:
    print("Reject H0,There is a relationship between 2 categorical variables")
else:
    print("Retain H0,There is no relationship between 2 categorical variables")

if p_value<=alpha:
    print("Reject H0,There is a relationship between 2 categorical variables")
else:
    print("Retain H0,There is no relationship between 2 categorical variables")
```

```
Degree of Freedom:- 1
chi-square statistic:- 54.04781105990782
critical_value: 3.841458820694124
p-value: 1.9562129693895258e-13
Significance level: 0.05
Degree of Freedom: 1
chi-square statistic: 54.04781105990782
critical_value: 3.841458820694124
p-value: 1.9562129693895258e-13
Reject H0,There is a relationship between 2 categorical variables
Reject H0,There is a relationship between 2 categorical variables
```

Therefore,variables ('tipovivi3','v2a1') have relationship between them. For good result we can use any one of them.

```
[19]: contingency_tab=pd.crosstab(train['v18q'],train['v18q1'])
Observed_Values=contingency_tab.values
import scipy.stats
b=scipy.stats.chi2_contingency(contingency_tab)
Expected_Values = b[3]
no_of_rows=len(contingency_tab.iloc[0:2,0])
no_of_columns=len(contingency_tab.iloc[0,0:2])
df=(no_of_rows-1)*(no_of_columns-1)
print("Degree of Freedom:-",df)
from scipy.stats import chi2
chi_square=sum([(o-e)**2./e for o,e in zip(Observed_Values,Expected_Values)])
chi_square_statistic=chi_square[0]+chi_square[1]
print("chi-square statistic:-",chi_square_statistic)
alpha=0.05
critical_value=chi2.ppf(q=1-alpha,df=df)
print('critical_value:',critical_value)
p_value=1-chi2.cdf(x=chi_square_statistic,df=df)
print('p-value:',p_value)
print('Significance level: ',alpha)
print('Degree of Freedom: ',df)
print('chi-square statistic:',chi_square_statistic)
print('critical_value:',critical_value)
print('p-value:',p_value)
if chi_square_statistic>=critical_value:
    print("Reject H0,There is a relationship between 2 categorical variables")
else:
    print("Retain H0,There is no relationship between 2 categorical variables")

if p_value<=alpha:
    print("Reject H0,There is a relationship between 2 categorical variables")
else:
    print("Retain H0,There is no relationship between 2 categorical variables")
```

```
Degree of Freedom:- 0
chi-square statistic:- 0.0
critical_value: nan
p-value: nan
Significance level: 0.05
Degree of Freedom: 0
chi-square statistic: 0.0
critical_value: nan
p-value: nan
Retain H0,There is no relationship between 2 categorical variables
Retain H0,There is no relationship between 2 categorical variables
```

Therefore,variables ('v18q','v18q1') have relationship between them. For good result we can use any one of them.

Conclusion : Therefore, there is bias in our dataset.

```
[21]: train.drop('r4t3',axis=1,inplace=True)
```

```

↳ -----
KeyError                                Traceback (most recent call↳
↳ last)

<ipython-input-21-5ca3a0904dd3> in <module>
----> 1 train.drop('r4t3',axis=1,inplace=True)

/usr/local/lib/python3.7/site-packages/pandas/core/frame.py in ↳
↳ drop(self, labels, axis, index, columns, level, inplace, errors)
    4172         level=level,
    4173         inplace=inplace,
-> 4174         errors=errors,
    4175     )
    4176

/usr/local/lib/python3.7/site-packages/pandas/core/generic.py in ↳
↳ drop(self, labels, axis, index, columns, level, inplace, errors)
    3887         for axis, labels in axes.items():
    3888             if labels is not None:
-> 3889                 obj = obj._drop_axis(labels, axis, level=level,↳
↳ errors=errors)
    3890
    3891         if inplace:

/usr/local/lib/python3.7/site-packages/pandas/core/generic.py in ↳
↳ _drop_axis(self, labels, axis, level, errors)
    3921         new_axis = axis.drop(labels, level=level,↳
↳ errors=errors)
    3922         else:
-> 3923             new_axis = axis.drop(labels, errors=errors)
    3924             result = self.reindex(**{axis_name: new_axis})
    3925

/usr/local/lib/python3.7/site-packages/pandas/core/indexes/base.py in ↳
↳ drop(self, labels, errors)
    5285         if mask.any():
```



```

5286             if errors != "ignore":
-> 5287                 raise KeyError(f"{labels[mask]} not found in axis")
5288                 indexer = indexer[~mask]
5289             return self.delete(indexer)

```

```

KeyError: "['r4t3'] not found in axis"

```

Check if there is a house without a family head

"parentesco1" =1 if household head

```

[22]: train.parentesco1.value_counts()

```

```

[22]: 0    6584
      1    2973
      Name: parentesco1, dtype: int64

```

```

[23]: pd.crosstab(train['edjefa'],train['edjefe'])

```

```

[23]: edjefe  0.0   1.0   2.0   3.0   4.0   5.0   6.0   7.0   8.0   9.0   ...  12.0  \
edjefa
0.0         435   123   194   307   137   222  1845   234   257   486   ...   113
1.0          69    0    0    0    0    0    0    0    0    0   ...    0
2.0          84    0    0    0    0    0    0    0    0    0   ...    0
3.0         152    0    0    0    0    0    0    0    0    0   ...    0
4.0         136    0    0    0    0    0    0    0    0    0   ...    0
5.0         176    0    0    0    0    0    0    0    0    0   ...    0
6.0         947    0    0    0    0    0    0    0    0    0   ...    0
7.0         179    0    0    0    0    0    0    0    0    0   ...    0
8.0         217    0    0    0    0    0    0    0    0    0   ...    0
9.0         237    0    0    0    0    0    0    0    0    0   ...    0
10.0          96    0    0    0    0    0    0    0    0    0   ...    0
11.0         399    0    0    0    0    0    0    0    0    0   ...    0
12.0          72    0    0    0    0    0    0    0    0    0   ...    0
13.0          52    0    0    0    0    0    0    0    0    0   ...    0
14.0         120    0    0    0    0    0    0    0    0    0   ...    0
15.0         188    0    0    0    0    0    0    0    0    0   ...    0
16.0         113    0    0    0    0    0    0    0    0    0   ...    0
17.0          76    0    0    0    0    0    0    0    0    0   ...    0
18.0           3    0    0    0    0    0    0    0    0    0   ...    0
19.0           4    0    0    0    0    0    0    0    0    0   ...    0
20.0           2    0    0    0    0    0    0    0    0    0   ...    0
21.0           5    0    0    0    0    0    0    0    0    0   ...    0

edjefe  13.0  14.0  15.0  16.0  17.0  18.0  19.0  20.0  21.0
edjefa

```

0.0	103	208	285	134	202	19	14	7	43
1.0	0	0	0	0	0	0	0	0	0
2.0	0	0	0	0	0	0	0	0	0
3.0	0	0	0	0	0	0	0	0	0
4.0	0	0	0	0	0	0	0	0	0
5.0	0	0	0	0	0	0	0	0	0
6.0	0	0	0	0	0	0	0	0	0
7.0	0	0	0	0	0	0	0	0	0
8.0	0	0	0	0	0	0	0	0	0
9.0	0	0	0	0	0	0	0	0	0
10.0	0	0	0	0	0	0	0	0	0
11.0	0	0	0	0	0	0	0	0	0
12.0	0	0	0	0	0	0	0	0	0
13.0	0	0	0	0	0	0	0	0	0
14.0	0	0	0	0	0	0	0	0	0
15.0	0	0	0	0	0	0	0	0	0
16.0	0	0	0	0	0	0	0	0	0
17.0	0	0	0	0	0	0	0	0	0
18.0	0	0	0	0	0	0	0	0	0
19.0	0	0	0	0	0	0	0	0	0
20.0	0	0	0	0	0	0	0	0	0
21.0	0	0	0	0	0	0	0	0	0

[22 rows x 22 columns]

Interpretation : Above cross tab shows 0 male head and 0 female head which implies that there are 435 families with no family head.

Count how many null values are existing in columns.

```
[24]: train.isna().sum().value_counts()
```

```
[24]: 0          135
      5           2
      7928         1
      6860         1
      7342         1
      dtype: int64
```

Lets Identify number of null values in Target variable

```
[25]: train['Target'].isna().sum()
```

```
[25]: 0
```

Interpretation : There are no null values in Target variable. Now lets proceed further and identify and fillna of other variable.

```
[26]: float_col=[]
      for i in train.columns:
          a=train[i].dtype
          if a == 'float64':
              float_col.append(i)
      print(float_col)
```

```
['v2a1', 'v18q1', 'rez_esc', 'dependency', 'edjefe', 'edjefa', 'meaneduc',
'overcrowding', 'SQBovercrowding', 'SQBdependency', 'SQBmeaned']
```

```
[27]: train[float_col].isna().sum()
```

```
[27]: v2a1          6860
      v18q1       7342
      rez_esc     7928
      dependency      0
      edjefe        0
      edjefa        0
      meaneduc       5
      overcrowding   0
      SQBovercrowding 0
      SQBdependency   0
      SQBmeaned      5
      dtype: int64
```

```
[28]: train['v18q1'].value_counts()
```

```
[28]: 1.0    1586
      2.0    444
      3.0    129
      4.0     37
      5.0     13
      6.0      6
      Name: v18q1, dtype: int64
```

```
[29]: pd.crosstab(train['tipovivi1'],train['v2a1'])
```

```
[29]: v2a1      0.0      12000.0      13000.0      14000.0      15000.0      16000.0  \
tipovivi1
0          29          3          4          3          3          2

v2a1      17000.0      20000.0      23000.0      25000.0      ...  570540.0  \
tipovivi1
0          4          22          5          21      ...          25

v2a1      600000.0      620000.0      684648.0      700000.0      770229.0      800000.0  \
tipovivi1
```

```

0          11          3          3          7          3          4

v2a1      855810.0   1000000.0  2353477.0
tipovivi1
0          11          7          2

[1 rows x 157 columns]

```

```
[30]: pd.crosstab(train['v18q1'],train['v18q'])
```

```

[30]: v18q      1
      v18q1
      1.0    1586
      2.0     444
      3.0    129
      4.0     37
      5.0     13
      6.0      6

```

Interpretation and action : *'v2a1', 'v18q1', 'rez_esc' have more than 50% null values, because for v18q1, there are families with their own house so they won't pay rent in that case it should be 0 and similar is for v18q1 there can be families with 0 tablets.*

Instead we can drop a column tipovivi3,v18q

- tipovivi3, =1 rented
- v18q, owns a tablet

as v2a1 alone can show both as v18q1 alone can show that if respondent owns a tablet or not

```
[31]: train['v2a1'].fillna(0,inplace=True)
      train['v18q1'].fillna(0,inplace=True)
```

```
[32]: train.drop(['tipovivi3', 'v18q','rez_esc','elimbasu5'],axis=1,inplace=True)
```

```
[33]: train['meaneduc'].fillna(np.mean(train['meaneduc']),inplace=True)
      train['SQBmeaned'].fillna(np.mean(train['SQBmeaned']),inplace=True)
      print(train.isna().sum().value_counts())
```

```

0      136
dtype: int64

```

```
[34]: int_col=[]
      for i in train.columns:
          a=train[i].dtype
          if a == 'int64':
              int_col.append(i)
      print(int_col)
```

```
['hacdor', 'rooms', 'hacapo', 'v14a', 'refrig', 'r4h1', 'r4h2', 'r4h3', 'r4m1',
'r4m2', 'r4m3', 'r4t1', 'r4t2', 'tamhog', 'tamviv', 'escolari', 'hhsiz',
'paredblolad', 'paredzocalo', 'paredpreb', 'pareddes', 'paredmad', 'paredzinc',
'paredfibras', 'paredother', 'pisomoscer', 'pisocemento', 'pisother',
'pisonatur', 'pisonotiene', 'pisomadera', 'techozinc', 'techoentrepiso',
'techocane', 'techootro', 'cielorazo', 'abastaguadentro', 'abastaguafuera',
'abastaguano', 'public', 'planpri', 'noelec', 'coopele', 'sanitario1',
'sanitario2', 'sanitario3', 'sanitario5', 'sanitario6', 'energcocinar1',
'energcocinar2', 'energcocinar3', 'energcocinar4', 'elimbasu1', 'elimbasu2',
'elimbasu3', 'elimbasu4', 'elimbasu6', 'epared1', 'epared2', 'epared3',
'etecho1', 'etecho2', 'etecho3', 'eviv1', 'eviv2', 'eviv3', 'dis', 'male',
'female', 'estadocivil1', 'estadocivil2', 'estadocivil3', 'estadocivil4',
'estadocivil5', 'estadocivil6', 'estadocivil7', 'parentesco1', 'parentesco2',
'parentesco3', 'parentesco4', 'parentesco5', 'parentesco6', 'parentesco7',
'parentesco8', 'parentesco9', 'parentesco10', 'parentesco11', 'parentesco12',
'hogar_nin', 'hogar_adul', 'hogar_mayor', 'hogar_total', 'instlevel1',
'instlevel2', 'instlevel3', 'instlevel4', 'instlevel5', 'instlevel6',
'instlevel7', 'instlevel8', 'instlevel9', 'bedrooms', 'tipovivi1', 'tipovivi2',
'tipovivi4', 'tipovivi5', 'computer', 'television', 'mobilephone',
'qmobilephone', 'lugar1', 'lugar2', 'lugar3', 'lugar4', 'lugar5', 'lugar6',
'area1', 'area2', 'age', 'SQBescolari', 'SQBage', 'SQBhogar_total', 'SQBedjefe',
'SQBhogar_nin', 'agesq', 'Target']
```

```
[35]: train[int_col].isna().sum().value_counts()
```

```
[35]: 0    126
      dtype: int64
```

Interpretation : Now there is no null value in our dataset.

```
[37]: train.Target.value_counts()
```

```
[37]: 4    5996
      2    1597
      3    1209
      1     755
      Name: Target, dtype: int64
```

Set the poverty level of the members and the head of the house same in a family.

Now for people below poverty level can be people paying less rent and don't own a house. and it also depends on whether a house is in urban area or rural area.

```
[38]: Poverty_level=train[train['v2a1'] !=0]
```

```
[39]: Poverty_level.shape
```

```
[39]: (2668, 136)
```

```
[40]: poverty_level=Poverty_level.groupby('area1')['v2a1'].apply(np.median)
```

```
[41]: poverty_level
```

```
[41]: area1
0      80000.0
1     140000.0
Name: v2a1, dtype: float64
```

- For rural area level if people paying rent less than 8000 is under poverty level.
- For Urban area level if people paying rent less than 140000 is under poverty level.

```
[42]: def povert(x):
      if x<8000:
          return('Below poverty level')

      elif x>140000:
          return('Above poverty level')
      elif x<140000:
          return('Below poverty level: Ur-ban ; Above poverty level : Rural ')
```

```
[43]: c=Poverty_level['v2a1'].apply(povert)
```

```
[44]: c.shape
```

```
[44]: (2668,)
```

```
[45]: pd.crosstab(c,Poverty_level['area1'])
```

```
[45]: area1              0      1
v2a1
Above poverty level          139  1103
Below poverty level: Ur-ban ; Above poverty lev...  306  1081
```

Interpretation : * *There are total 1242 people above poverty level independent of area whether rural or Urban* * *Remaining 1111 people level depends on their area*

Rural :

Above poverty level= 445

Urban :

Above poverty level =1103

Below poverty level=1081

```
[46]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.model_selection import train_test_split
```

```
[47]: X_data=train.drop('Target',axis=1)
      Y_data=train.Target
```

```
[48]: X_data_col=X_data.columns
```

Applying Standard Scalling to dataset

```
[49]: from sklearn.preprocessing import StandardScaler
      SS=StandardScaler()
      X_data_1=SS.fit_transform(X_data)
      X_data_1=pd.DataFrame(X_data_1,columns=X_data_col)
```

Now we will proceed to model fitting

```
[50]: X_train,X_test,Y_train,Y_test=train_test_split(X_data_1,Y_data,test_size=0.
      ↪25,stratify=Y_data,random_state=0)
```

Lets identify best parameters for our model using GridSearchCv

```
[51]: from sklearn.pipeline import Pipeline
      from sklearn.model_selection import GridSearchCV

      rfc=RandomForestClassifier(random_state=0)
      parameters={'n_estimators':[10,50,100,300], 'max_depth':[3,5,10,15]}
      grid=zip([rfc],[parameters])

      best_=None

      for i, j in grid:
          a=GridSearchCV(i,param_grid=j,cv=3,n_jobs=1)
          a.fit(X_train,Y_train)
          if best_ is None:
              best_=a
          elif a.best_score_>best_.best_score_:
              best_=a

      print ("Best CV Score",best_.best_score_)
      print ("Model Parameters",best_.best_params_)
      print("Best Estimator",best_.best_estimator_)
```

Best CV Score 0.8507046183898423

Model Parameters {'max_depth': 15, 'n_estimators': 300}

Best Estimator RandomForestClassifier(max_depth=15, n_estimators=300,
random_state=0)

```
[52]: RFC=best_.best_estimator_
      Model=RFC.fit(X_train,Y_train)
```

```
pred=Model.predict(X_test)
```

```
[53]: print('Model Score of train data : {}'.format(Model.score(X_train,Y_train)))  
      print('Model Score of test data : {}'.format(Model.score(X_test,Y_test)))
```

Model Score of train data : 0.9831170643225896

Model Score of test data : 0.8824267782426778

```
[54]: Important_features=pd.DataFrame(Model.  
      ↪feature_importances_,X_data_col,columns=['feature_importance'])
```

```
[55]: Top50Features=Important_features.  
      ↪sort_values(by='feature_importance',ascending=False).head(50).index
```

```
[56]: Top50Features
```

```
[56]: Index(['SQBmeaned', 'meaneduc', 'SQBdependency', 'dependency', 'overcrowding',  
            'SQBovercrowding', 'qmobilephone', 'SQBhogar_nin', 'SQBedjefe',  
            'edjefe', 'hogar_nin', 'rooms', 'cielorazo', 'r4t1', 'v2a1', 'edjefa',  
            'agesq', 'r4m3', 'r4h2', 'SQBage', 'age', 'escolari', 'r4t2', 'r4h3',  
            'hogar_adul', 'SQBescolari', 'eviv3', 'bedrooms', 'r4m1', 'epared3',  
            'r4m2', 'tamviv', 'paredblolad', 'v18q1', 'SQBhogar_total', 'tamhog',  
            'hhszise', 'hogar_total', 'pisomoscer', 'etecho3', 'r4h1', 'lugar1',  
            'eviv2', 'tipovivi1', 'energcocinar2', 'energcocinar3', 'epared2',  
            'television', 'area2', 'area1'],  
           dtype='object')
```

```
[57]: for i in Top50Features:  
      if i not in X_data_col:  
          print(i)
```

```
[58]: X_data_Top50=X_data[Top50Features]
```

```
[59]: X_train,X_test,Y_train,Y_test=train_test_split(X_data_Top50,Y_data,test_size=0.  
      ↪25,stratify=Y_data,random_state=0)
```

```
[60]: Model_1=RFC.fit(X_train,Y_train)  
      pred=Model_1.predict(X_test)
```

```
[61]: from sklearn.metrics import confusion_matrix,f1_score,accuracy_score
```

```
[62]: confusion_matrix(Y_test,pred)
```

```
[62]: array([[ 143,   17,    0,   29],  
            [   8,  324,    4,   63],  
            [   1,   12,  214,   75],  
            [   2,   10,    3, 1485]])
```



```
[63]: f1_score(Y_test,pred,average='weighted')
```

```
[63]: 0.9026906492316511
```

```
[64]: accuracy_score(Y_test,pred)
```

```
[64]: 0.906276150627615
```

Lets apply cleaning on test data and then find prediction for that.

```
[65]: # lets drop Id variable.
test.drop('r4t3',axis=1,inplace=True)
test.drop(['Id','idhogar'],axis=1,inplace=True)
test['dependency']=test['dependency'].apply(map)
test['edjefe']=test['edjefe'].apply(map)
test['edjefa']=test['edjefa'].apply(map)
```

```
[66]: test['v2a1'].fillna(0,inplace=True)
test['v18q1'].fillna(0,inplace=True)
```

```
[67]: test.drop(['tipovivi3', 'v18q','rez_esc','elimbasu5'],axis=1,inplace=True)
```

```
[68]: train['meaneduc'].fillna(np.mean(train['meaneduc']),inplace=True)
train['SQBmeaned'].fillna(np.mean(train['SQBmeaned']),inplace=True)
```

```
[69]: test_data=test[Top50Features]
```

```
[70]: test_data.isna().sum().value_counts()
```

```
[70]: 0      48
31      2
dtype: int64
```

```
[71]: test_data.SQBmeaned.fillna(np.mean(test_data['SQBmeaned']),inplace=True)
```

/usr/local/lib/python3.7/site-packages/pandas/core/series.py:4536:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
downcast=downcast,

```
[72]: test_data.meaneduc.fillna(np.mean(test_data['meaneduc']),inplace=True)
```

```
[73]: Test_data_1=SS.fit_transform(test_data)
X_data_1=pd.DataFrame(Test_data_1)
```

```
[74]: test_prediction=Model_1.predict(test_data)
```

```
[75]: test_prediction
```

```
[75]: array([4, 4, 4, ..., 4, 4, 4])
```

Interpretation : *Above is our prediction for test data.*

Conclusion :

Using RandomForest Classifier we can predict test_data with accuracy of 90%.

```
[ ]:
```